

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [3]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
```

```
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```
In [4]: # using SQLite Table to read data.
```

```

con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500
000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 11500
0""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative
rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (115000, 10)

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summa
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality D Food
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertise
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all

```
In [5]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [6]: print(display.shape)
display.head()
```

```
(80668, 7)
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [7]: display[display['UserId']=='AZY10LLTJ71NX']
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

```
In [8]: display['COUNT(*)'].sum()
```

```
393063
```

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [9]: display= pd.read_sql_query("""
SELECT *
```

```
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Sum
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [10]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False,
, kind='quicksort', na_position='last')
```

```
In [11]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

(99724, 10)

```
In [12]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

86.71652173913044

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from

calculations

```
In [13]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5	1224892800	Bought This for Son at College
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4	1212883200	Pure co taste wi crunchy almond inside

```
In [14]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [15]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(99722, 10)
```

```
1    83711
```

```
0    16011
```

```
Name: Score, dtype: int64
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [16]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)
```

```

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont t ake any chances till they know what is going on with the china imports.

=====

IF YOU LIKE SALMON YOU WILL LOVE THESE OMAHA STEAKS SALMON VERY VERY GOOD

=====

OK...I thought I'd put a bit of punch to hubby's sandwich, instead of the ho-hum Best Foods Mayo---oh0ooo0h--FAILURE!
One bite and he said---Please! DO NOT EVER SERVE THIS TO ME AGAIN!
I guess it was that-bad!
I'll see if my neighbo r will be able to use it w/her family.
If you are a BEST FOODS lover---walk away---do NOT purchase this product!

=====

These people from Bavaria really know how to make this stuff. The Landjagers are super (you have to let them dry for them to develop their full, intended flavor), and it is worth any sausage fan's time to check out their complete offering of German style sausages and hams. Due to the perishability of some of their products their S&H charges appear outrageous but I guess that sending frozen or refrigerated foods costs money. Personally, I recommend their coarse grind liverwurst but that is a matter of personal taste.

=====

```
In [17]: final['Text'].values[0]
```

'My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicke n products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.'

```
In [18]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_150)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont t ake any chances till they know what is going on with the china imports.

```
In [19]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-ta
         gs-from-an-element
         from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)
```

```
soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont t ake any chances till they know what is going on with the china imports.

=====

IF YOU LIKE SALMON YOU WILL LOVE THESE OMAHA STEAKS SALMON VERY VERY GOOD

=====

OK...I thought I'd put a bit of punch to hubby's sandwich, instead of the ho-hum Best Foods Mayo---oh0ooo0h--FAILURE!One bite and he said---Please! DO NOT EVER SERVE THIS TO ME AGAIN!I guess it was that-bad!I'll see if my neighbor will be able to use it w/her family.If you are a BEST FOODS lover---walk away---do NOT purchase this product!

=====

These people from Bavaria really know how to make this stuff. The Landjagers are super (you have to let them dry for them to develop their full, intended flavor), and it is worth any sausage fan's time to check out their complete offering of German style sausages and hams. Due to the perishability of some of their products their S&H charges appear outrageous but I guess that sending frozen or refrigerated foods costs money. Personally, I recommend their coarse grind liverwurst but that is a matter of personal taste.

```
In [20]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
```

```
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase
```

```
In [21]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

```
OK...I thought I would put a bit of punch to hubby is sandwich, instead of the ho-hum Best Foods Mayo---oh0ooo0h--FAILURE!<br />One bite and he said---Please! DO NOT EVER SERVE THIS TO ME AGAIN!<br />I guess it was that-bad!<br />I will see if my neighbor will be able to use it w/her family.<br />If you are a BEST FOODS lover---walk away---do NOT purchase this product!
=====
```

```
In [22]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.
```

```
In [23]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

```
OK I thought I would put a bit of punch to hubby is sandwich instead of the ho hum Best Foods Mayo oh0ooo0h FAILURE br One bite and he said Please DO NOT EVER SERVE THIS TO ME AGAIN br I guess it was that bad br I will see if my neighbor will be able to use it w her family br If you are a BEST FOODS lover walk away do NOT purchase this product
```

```
In [24]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
```

```

# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
, 'you', "you're", "you've",\
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'his', 'himself', \
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'th
ey', 'them', 'their',\
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "tha
t'll", 'these', 'those', \
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'ha
d', 'having', 'do', 'does', \
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
'until', 'while', 'of', \
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
'during', 'before', 'after',\
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'ove
r', 'under', 'again', 'further',\
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any'
, 'both', 'each', 'few', 'more',\
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
'very', \
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no
w', 'd', 'll', 'm', 'o', 're', \
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
'doesn', "doesn't", 'hadn',\
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'might
n', "mightn't", 'mustn',\

```

```
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [25]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|██████████| 99722/99722 [00:47<00:00, 2094.79it/s]

```
In [26]: preprocessed_reviews[1500]
```

'ok thought would put bit punch hubby sandwich instead ho hum best foods mayo ohoooooh failure one bite said please not ever serve guess bad see neighbor able use w family best foods lover walk away not purchase product'

[3.2] Preprocessing Review Summary

```
In [27]: ## Similarly you can do preprocessing for review summary also.
```



```
In [28]: from sklearn.cross_validation import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.cross_validation import cross_val_score
from collections import Counter
import collections
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve, confusion_matrix, auc
from sklearn import cross_validation
from scipy.sparse import csr_matrix, hstack
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

/usr/local/lib/python3.5/dist-packages/sklearn/cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that at the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)

```
In [29]: X_1, X_test, y_1, y_test = cross_validation.train_test_split(preprocessed_reviews, final[
'Score'], test_size=0.2, random_state=0)
X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_1, y_1, test_size=0.25)
print(np.asarray(X_1).shape, np.asarray(X_test).shape, np.asarray(X_tr).shape, np.asarray(X
_test).shape, np.asarray(X_cv).shape)
```

(79777,) (19945,) (59832,) (19945,) (19945,)

[4] Featurization

[4.1] BAG OF WORDS

```
In [30]: #BoW
```

```

count_vect = CountVectorizer() #in scikit-learn
BOW_Train = count_vect.fit_transform(X_tr)
BOW_test = count_vect.transform(X_test)
BOW_CV = count_vect.transform(X_cv)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)
print("the type of count vectorizer ",type(BOW_Train))
print("the shape of out text BOW vectorizer ",BOW_Train.get_shape())
print("the number of unique words ", BOW_Train.get_shape()[1])

```

```

some feature names  ['aa', 'aaa', 'aaaa', 'aaaaa', 'aaaaaaaaaaaa', 'aaaand', 'aaah', 'aaahs', 'aadp', 'aaf']
=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (59832, 45877)
the number of unique words  45877

```

[4.2] Bi-Grams and n-Grams.

```

In [0]: #bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature\_extraction.text.CountVectorizer.html

# you can choose these numebtrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())

```

```
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>  
the shape of out text BOW vectorizer (4986, 3144)  
the number of unique words including both unigrams and bigrams 3144
```

[4.3] TF-IDF

```
In [31]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)  
TFIDF_Train = tf_idf_vect.fit_transform(X_tr)  
TFIDF_Test = tf_idf_vect.transform(X_test)  
TFIDF_Validation = tf_idf_vect.transform(X_cv)  
print("the type of count vectorizer ", type(TFIDF_Train))  
print("the shape of out text TFIDF vectorizer ", TFIDF_Train.get_shape())  
print("the number of unique words including both unigrams and bigrams ", TFIDF_Train.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>  
the shape of out text TFIDF vectorizer (59832, 35214)  
the number of unique words including both unigrams and bigrams 35214
```

[4.4] Word2Vec

```
In [32]: # Train your own Word2Vec model using your own text corpus  
i=0  
list_of_sentence=[]  
list_of_sentence_cv=[]  
list_of_sentence_test=[]  
for sentence in X_tr:
```

```
list_of_sentence.append(sentence.split())
for sentence in X_cv:
    list_of_sentence_cv.append(sentence.split())
for sentence in X_test:
    list_of_sentence_test.append(sentence.split())
```

```
In [33]: # Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occured atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
```

```

print('='*50)
print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin'
, binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, to
train your own w2v ")

```

```

[('awesome', 0.821196973323822), ('fantastic', 0.8178302645683289), ('good', 0.8044040203094482), ('wonderful', 0.8031048774
719238), ('excellent', 0.7855092883110046), ('terrific', 0.7614560127258301), ('perfect', 0.7441087365150452), ('fabulous',
0.6897355318069458), ('amazing', 0.6886729001998901), ('nice', 0.6685606241226196)]
=====
[('greatest', 0.7765495181083679), ('nastiest', 0.7263230085372925), ('best', 0.6905359029769897), ('softest', 0.66272819042
20581), ('disgusting', 0.635009229183197), ('tastiest', 0.621982216835022), ('ive', 0.6203436851501465), ('awful', 0.6192965
507507324), ('horrible', 0.6023806929588318), ('terrible', 0.5971200466156006)]

```

```

In [34]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occured minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

```

```

number of words that occured minimum 5 times 14581
sample words ['requires', 'pocket', 'mmmmmmmm', 'reads', 'mentality', 'tremendous', 'stage', 'tinge', 'agent', 'bricks', 'q
uintessential', 'duh', 'cafeteria', 'behalf', 'haribos', 'bummed', 'rushed', 'amounts', 'larva', 'yg', 'fishy', 'testing',
'kmr', 'fleas', 'brittle', 'lara', 'continuously', 'sourced', 'walls', 'knife', 'breadsticks', 'marketing', 'applaud', 'rabb
it', 'hum', 'frequently', 'poppy', 'pearl', 'caff', 'marker', 'quickest', 'beeswax', 'hello', 'tear', 'brie', 'apiece', 'hes
itated', 'unpleasant', 'feels', 'persons']

```

[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```
In [35]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to c
    hange this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to c
    hange this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
```

```

        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
        sent_vectors_cv.append(sent_vec)
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to c
hange this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
        sent_vectors_test.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

```

100%|██████████| 59832/59832 [10:31<00:00, 94.75it/s]
100%|██████████| 19945/19945 [03:32<00:00, 94.01it/s]
100%|██████████| 19945/19945 [03:29<00:00, 95.05it/s]

```

```

59832
50

```

[4.4.1.2] TFIDF weighted W2v

```
In [36]: model = TfidfVectorizer()
```

```

tf_idf_matrix = model.fit_transform(X_tr)
tf_idf_matrix_cv = model.transform(X_cv)
tf_idf_matrix_test = model.transform(X_test)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

In [37]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names()# tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #         tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
tfidf_sent_vectors_cv = []; # the tfidf-w2v for each sentence/review is stored in this l

```



```

ist
row=0;
for sent in tqdm(list_of_sentence_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
#             # to reduce the computation we are
#             # dictionary[word] = idf value of word in whole corpus
#             # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_cv.append(sent_vec)
    row += 1
tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stored in this
list
row=0;
for sent in tqdm(list_of_sentence_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
#             # to reduce the computation we are

```

```

        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_test.append(sent_vec)
    row += 1

```

```

100%|██████████| 59832/59832 [58:20<00:00, 17.09it/s]
100%|██████████| 19945/19945 [23:33<00:00, 14.11it/s]
100%|██████████| 19945/19945 [23:22<00:00, 11.59it/s]

```

[5] Assignment 8: Decision Trees

1. Apply Decision Trees on these feature sets

- **SET 1:** Review text, preprocessed one converted into vectors using (BOW)
- **SET 2:** Review text, preprocessed one converted into vectors using (TFIDF)
- **SET 3:** Review text, preprocessed one converted into vectors using (AVG W2v)
- **SET 4:** Review text, preprocessed one converted into vectors using (TFIDF W2v)

2. The hyper parameter tuning (best `depth` in range [1, 5, 10, 50, 100, 500, 100], and the best `min_samples_split` in range [5, 10, 100, 500])

- Find the best hyper parameter which will give the maximum [AUC](#) value
- Find the best hyper parameter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

3. Graphviz

- Visualize your decision tree with Graphviz. It helps you to understand how a decision is being made, given a new vector.
- Since feature names are not obtained from word2vec related models, visualize only BOW & TFIDF decision trees using Graphviz
- Make sure to print the words in each node of the decision tree instead of printing its index.
- Just for visualization purpose, limit max_depth to 2 or 3 and either embed the generated images of graphviz in your notebook, or directly upload them as .png files.

4. Feature importance


- Find the top 20 important features from both feature sets **Set 1** and **Set 2** using `feature_importances_`` method of [Decision Tree Classifier](#) and print their corresponding feature names


5. Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

6. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure.

 Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.

 Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please visualize your confusion

matrices using [seaborn heatmaps](#).



7. Conclusion

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)



Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this [link](#).

Applying Decision Trees

[5.1] Applying Decision Trees on BOW, SET 1

```
In [39]: # Please write all the code with proper documentation
depth  =[5,10,50,100,500,1000]
split  = [5,10,50,100,500,1000]
BOW_Train_score = []
BOW_val_score = []
for i in depth:
    bow_tr = []
    bow_val = []
    for j in split:
```

```

model= DecisionTreeClassifier(max_depth=i,min_samples_split=j)
model.fit(BOW_Train,y_tr)
train_data = model.predict_log_proba(BOW_Train)[: ,1]
val_data = model.predict_log_proba(BOW_CV)[: ,1]
train_data[np.isinf(train_data)]=0
train_data[np.isnan(train_data)]=0
val_data[np.isinf(val_data)]=0
val_data[np.isnan(val_data)]=0

bow_tr.append(roc_auc_score(np.asarray(y_tr),np.asarray(train_data)))
bow_val.append(roc_auc_score(np.asarray(y_cv),np.asarray(val_data)))
print(str(bow_tr)+" "+str(bow_val))
BOW_Train_score.append(bow_tr)
BOW_val_score.append(bow_val)

```

```

[0.7102287347535805, 0.7102735983760587, 0.7116340566246048, 0.7116126100929957, 0.7119474016401816, 0.711416872395976]
[0.7047000868276412, 0.7048895961099575, 0.7059986978786693, 0.7061076082042369, 0.7055622866978029, 0.704717014991516]
[0.7403161081822004, 0.7452819978849052, 0.7694547063822865, 0.7699419125471645, 0.7686721490798217, 0.7671467754782769]
[0.7343338229516794, 0.7355505244193025, 0.755384714199802, 0.758150933818743, 0.7588061646598742, 0.7577917854540458]
[0.44464578174444674, 0.5019907574925818, 0.6654513962602209, 0.7194987485904969, 0.8516081704897813, 0.8484215966448054]
[0.4639402155542752, 0.49117592649756414, 0.6421001125929116, 0.6966191203088548, 0.8328267471074928, 0.8294527908245722]
[0.3793817806599712, 0.44107315884949566, 0.619649044439493, 0.6946711172081551, 0.8452728965730679, 0.8449037882791051]
[0.40001578434424745, 0.4281552044452937, 0.5939283781209594, 0.6663231236188929, 0.825772963608792, 0.8261255038488797]
[0.5408999886484097, 0.5847263642584637, 0.6996546438337298, 0.7495309672344008, 0.8644896067819969, 0.8604087831249653]
[0.5098096464174909, 0.5290246063464099, 0.6379112398168891, 0.6874191077561675, 0.812952917744558, 0.8133525397268352]
[0.539618889246679, 0.5844948652945311, 0.7009060032378125, 0.7364635312313395, 0.8645327566024067, 0.8598176863204486]
[0.5075835149635803, 0.5298639389892477, 0.6381154776099508, 0.6760563320901383, 0.811267525186395, 0.8138054574713931]


```

```

In [55]: df = pd.DataFrame(BOW_Train_score,index = depth,columns=split)
df2 = pd.DataFrame(BOW_val_score,index = depth,columns=split)
sns.heatmap(df, annot= True, fmt="f", cmap="YlGnBu")


```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7f1a7f28>
```



```
In [56]: sns.heatmap(df2, annot=True, fmt="f", cmap="YlGnBu")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7f0944e0>
```




```
In [49]: best_depth = 50
         best_split = 500
```

Testing on test data

```
In [81]: model= DecisionTreeClassifier(max_depth=best_depth,min_samples_split=best_split)
         model.fit(BOW_Train,y_tr)
         train_pred = model.predict_log_proba(BOW_Train)[:,-1]
         test_pred = model.predict_log_proba(BOW_test)[:,-1]
         train_pred[np.isinf(train_pred)]=0
         train_pred[np.isnan(train_pred)]=0
         test_pred[np.isinf(test_pred)]=0
         test_pred[np.isnan(test_pred)]=0
         train_fpr, train_tpr, thresholds = roc_curve(y_tr, train_pred)
         test_fpr, test_tpr, thresholds = roc_curve(y_test, test_pred)

         plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
```

```
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```




Confusion Matrix

```
In [57]: ytrain = model.predict(BOW_Train)
ytest = model.predict(BOW_test)
ctrain = confusion_matrix(y_tr, ytrain)
ctest = confusion_matrix(y_test, ytest)
```




```
class_label=["No","Yes"]
df = pd.DataFrame(ctest, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7f0377b8>



```
In [58]: df = pd.DataFrame(ctrain, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7effbac8>



[5.1.1] Top 20 important features from SET 1

```
In [87]: # Please write all the code with proper documentation
a = model.feature_importances_
b = []
for i in range(45884):
    b.append((i,a[i]))
b = sorted(b,key= lambda x: x[1],reverse=True)
b = b[0:20]
print(b)
print(" So the top 10 features of positive class are--")
for i in range(0,20):
    print("feature name : %s , value : %f"%(count_vect.get_feature_names()[b[i][0]],b[i][1]))
```

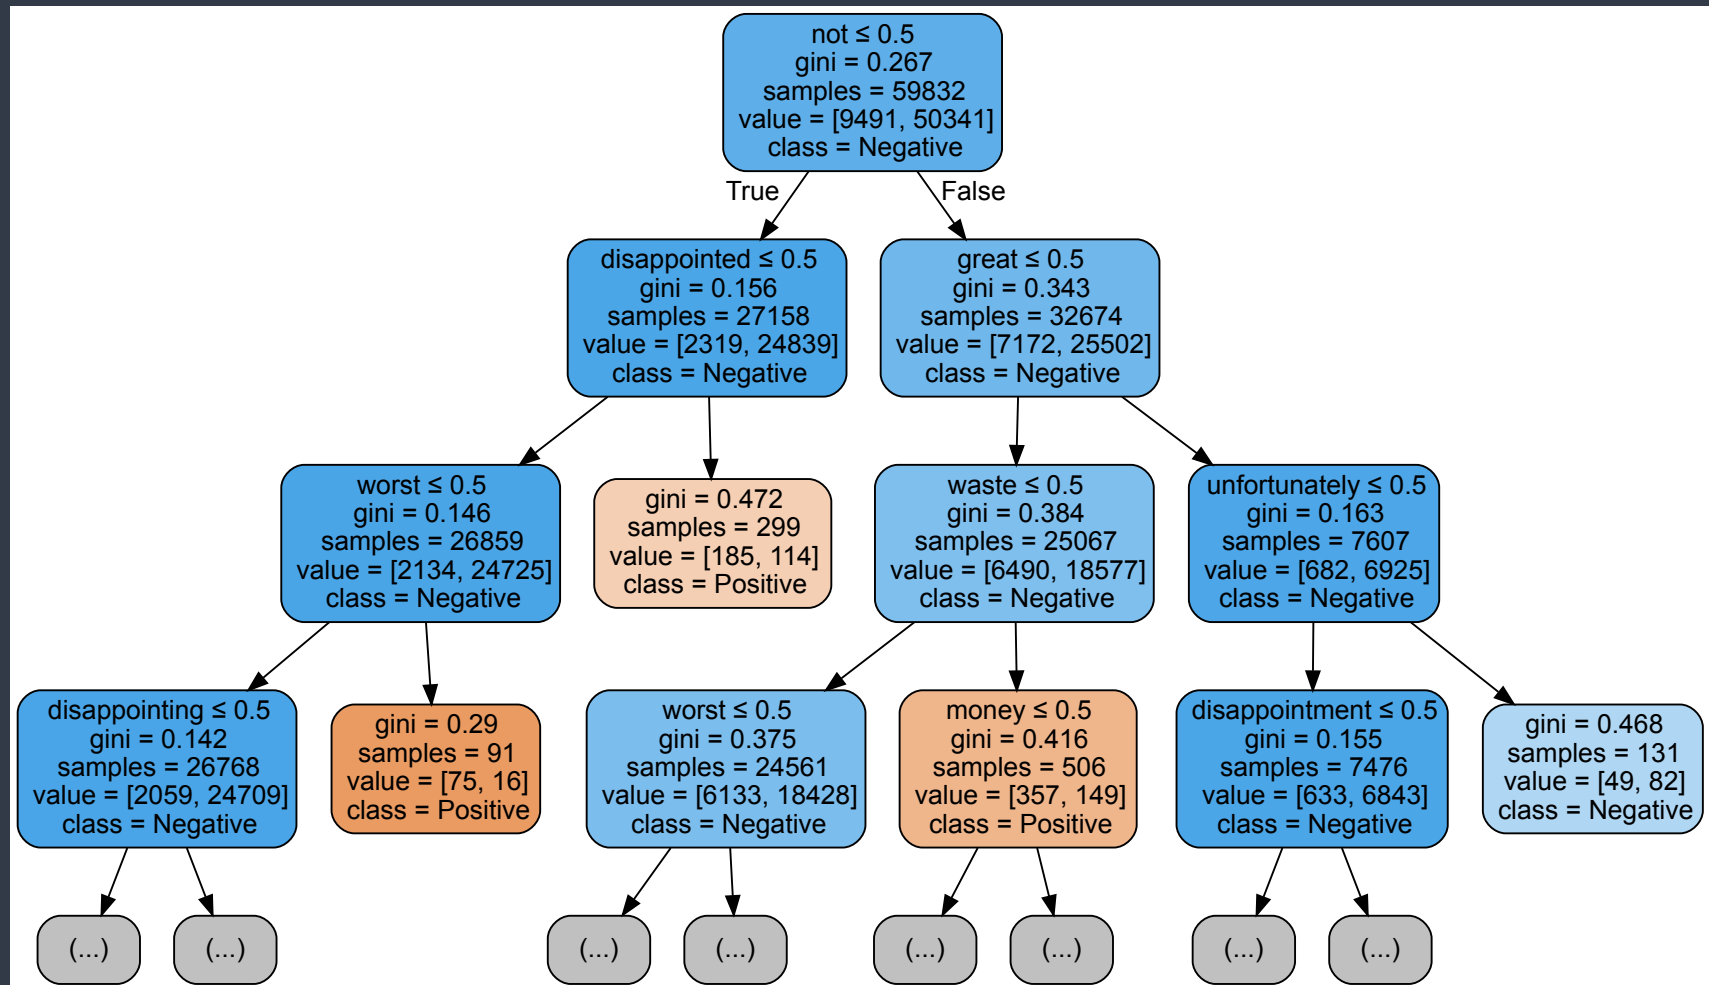
```
[(27089, 0.10164339719737056), (17462, 0.05800728633271654), (11340, 0.047599342005398046), (45174, 0.04308265280904299), (44200, 0.03522560134606249), (2738, 0.03047867970618844), (33890, 0.024437884554222595), (10484, 0.022705896772071504), (3704, 0.02209907236594975), (23529, 0.020222590452375826), (17089, 0.018342513508168946), (19175, 0.01686457125835333), (11343,
```

```
0.016781086216510015), (29557, 0.016651651962431145), (2867, 0.015948456210760995), (23547, 0.015482310894936123), (11345, 0.014337430343554003), (41013, 0.012920022239641756), (40463, 0.012870428423964315), (14496, 0.011777386920358375)]  
So the top 10 features of positive class are--  
feature name : not , value : 0.101643  
feature name : great , value : 0.058007  
feature name : disappointed , value : 0.047599  
feature name : worst , value : 0.043083  
feature name : waste , value : 0.035226  
feature name : awful , value : 0.030479  
feature name : return , value : 0.024438  
feature name : delicious , value : 0.022706  
feature name : best , value : 0.022099  
feature name : love , value : 0.020223  
feature name : good , value : 0.018343  
feature name : horrible , value : 0.016865  
feature name : disappointing , value : 0.016781  
feature name : perfect , value : 0.016652  
feature name : bad , value : 0.015948  
feature name : loves , value : 0.015482  
feature name : disappointment , value : 0.014337  
feature name : threw , value : 0.012920  
feature name : terrible , value : 0.012870  
feature name : favorite , value : 0.011777
```

[5.1.2] Graphviz visualization of Decision Tree on BOW, SET 1

```
In [152]: # Please write all the code with proper documentation  
f_name = count_vect.get_feature_names()  
from sklearn.tree import DecisionTreeClassifier, export_graphviz  
from sklearn import tree  
from sklearn.datasets import load_wine  
from IPython.display import SVG  
from graphviz import Source  
from IPython.display import display  
dot_data = Source(export_graphviz(model, feature_names = f_name , class_names=['Positive',
```

```
'Negative'] ,max_depth = 3,out_file=None,filled=True,rounded=True,special_characters=True))
display(SVG(dot_data.pipe(format='svg'))))
```



[5.2] Applying Decision Trees on TFIDF, SET 2

```

In [59]: # Please write all the code with proper documentation
depth = [5,10,50,100,500,1000]
split = [5,10,50,100,500,1000]
TFIDF_Train_score = []
TFIDF_val_score = []
for i in depth:
    tfidf_tr = []
    tfidf_val = []
    for j in split:
        model= DecisionTreeClassifier(max_depth=i,min_samples_split=j)
        model.fit(TFIDF_Train,y_tr)
        train_data = model.predict_log_proba(TFIDF_Train)[:,-1]
        val_data = model.predict_log_proba(TFIDF_Validation)[:,-1]
        train_data[np.isinf(train_data)]=0
        train_data[np.isnan(train_data)]=0
        val_data[np.isinf(val_data)]=0
        val_data[np.isnan(val_data)]=0

        tfidf_tr.append(roc_auc_score(np.asarray(y_tr),np.asarray(train_data)))
        tfidf_val.append(roc_auc_score(np.asarray(y_cv),np.asarray(val_data)))
    print(str(tfidf_tr)+" "+str(tfidf_val))
    TFIDF_Train_score.append(bow_tr)
    TFIDF_val_score.append(bow_val)

```

```


[0.7030555483538294, 0.7030555483538294, 0.7032490545594837, 0.7031776885849219, 0.7033587514603327, 0.701079075328427]
[0.6931734095078099, 0.6934065452438586, 0.6937574632397634, 0.6937053222953762, 0.6945014867455501, 0.6915318908643566]
[0.7357874801519385, 0.7377272922165289, 0.755621945033003, 0.7457881654640202, 0.7624772597670123, 0.7592228227000577]
[0.713724246277857, 0.7155744570119589, 0.7361301788449973, 0.7291578375034113, 0.7514126355701214, 0.7483018677841295]
[0.331225935014713, 0.3847329877590278, 0.5546126147688991, 0.6600912318838492, 0.843800229169362, 0.8507481894129951]
[0.39811767616325355, 0.4282173813070837, 0.5583443328210097, 0.6539823638411046, 0.8128070990725236, 0.8217479466604647]
[0.34177344112793473, 0.4131889980086079, 0.5774101672341001, 0.6661335145329371, 0.8366556440595507, 0.8426209976615892]

```

```
[0.3905275439098424, 0.4291229418395582, 0.545216500492649, 0.6328996601095985, 0.7974827206248423, 0.8034309767911665]  
[0.5296714873142535, 0.5774697735205349, 0.692224029959848, 0.7433460121287088, 0.855605087775264, 0.8591841901362572]  
[0.5091840559022779, 0.5346085449631829, 0.6152620164484195, 0.6684658148974294, 0.7829985638464718, 0.7947000402683166]  
[0.5308861596643136, 0.5786962942756381, 0.6876262373870117, 0.7449251815805648, 0.8553172368701156, 0.8572568464513924]  
[0.5110471254419652, 0.5307493635303671, 0.6151811425350298, 0.6677488196294699, 0.7801156251333539, 0.7938172828066166]
```


```
In [60]: df = pd.DataFrame(TFIDF_Train_score,index = depth,columns=split)  
df2 = pd.DataFrame(TFIDF_val_score,index = depth,columns=split)  
sns.heatmap(df, annot=True, fmt="f", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7ef53c88>



```
In [61]: sns.heatmap(df2, annot=True, fmt="f", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7effb390>




```
In [155]: best_depth = 50
          best_split = 500
```

Testing on test data

```
In [156]: model = DecisionTreeClassifier(max_depth=best_depth, min_samples_split=best_split)
          model.fit(TFIDF_Train, y_tr)
          train_pred = model.predict_log_proba(TFIDF_Train)[:, 1]
          test_pred = model.predict_log_proba(TFIDF_Test)[:, 1]
          train_pred[np.isinf(train_pred)] = 0
          train_pred[np.isnan(train_pred)] = 0
          test_pred[np.isinf(test_pred)] = 0
          test_pred[np.isnan(test_pred)] = 0
          train_fpr, train_tpr, thresholds = roc_curve(y_tr, train_pred)
          test_fpr, test_tpr, thresholds = roc_curve(y_test, test_pred)

          plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
```

```
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.plot([0.0,1.0],[0.0,1.0])
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```




Confusion Matrix

```
In [64]: ytrain = model.predict(TFIDF_Train)
ytest = model.predict(TFIDF_Test)
ctrain = confusion_matrix(y_tr,ytrain)
ctest = confusion_matrix(y_test,ytest)
class_label=["No","Yes"]
```




```
df = pd.DataFrame(ctest, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7eeabef0>




```
In [65]: df = pd.DataFrame(ctrain, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7ee7a4e0>



```
In [63]: df = pd.DataFrame(ctrain, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7eee3c50>



[5.2.1] Top 20 important features from SET 2

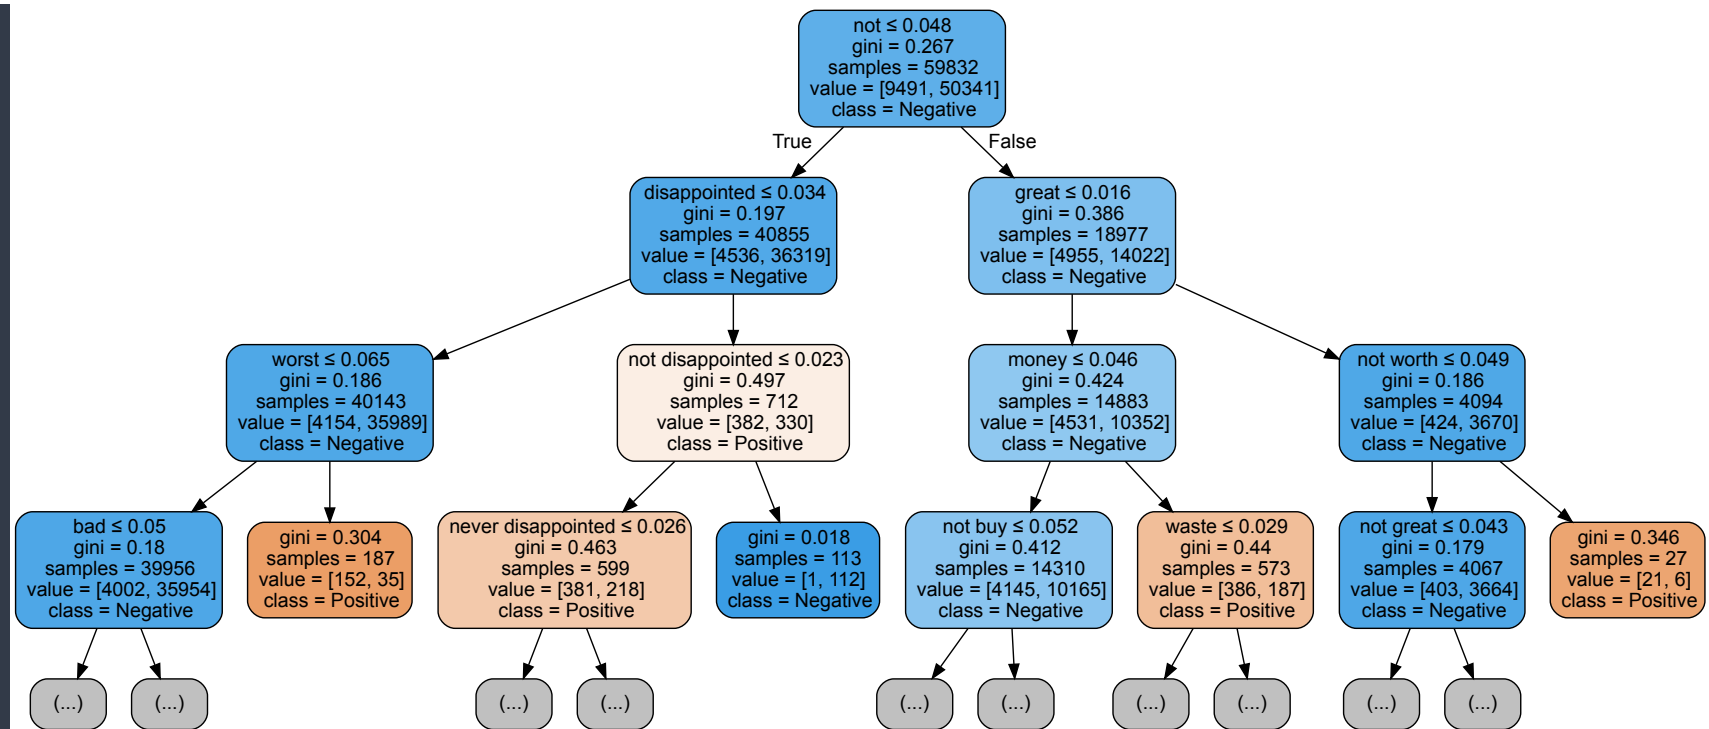
```
In [159]: # Please write all the code with proper documentation
a = model.feature_importances_
b = []
for i in range(35256):
    b.append((i,a[i]))
b = sorted(b,key= lambda x: x[1],reverse=True)
b = b[0:20]
print(b)
print(" So the top 10 features of positive class are--")
for i in range(0,20):
    print("feature name : %s , value : %f"%(tf_idf_vect.get_feature_names()[b[i][0]],b[i][1]))

[(20545, 0.10081024283327579), (13532, 0.05376215851996582), (7842, 0.045071790360250476), (34797, 0.03970912535227808), (20627, 0.02441228428845962), (19400, 0.024164473756463106), (1688, 0.021870228444515364), (1776, 0.020905197552283563), (14733, 0.019544588115575544), (7857, 0.01865231105079609), (21288, 0.017035360975440044), (33820, 0.01695065299824854), (26066, 0.01688956397140815), (7373, 0.016376347191560267), (31151, 0.01532519797597335), (31580, 0.014951342241692351), (2356, 0.014249703296028038), (21077, 0.014124863010970037), (12954, 0.013402922194049491), (17671, 0.013281680876820037)]
So the top 10 features of positive class are--
feature name : not , value : 0.100810
feature name : great , value : 0.053762
feature name : disappointed , value : 0.045072
feature name : worst , value : 0.039709
feature name : not buy , value : 0.024412
feature name : money , value : 0.024164
feature name : awful , value : 0.021870
feature name : bad , value : 0.020905
feature name : horrible , value : 0.019545
feature name : disappointing , value : 0.018652
feature name : not worth , value : 0.017035
feature name : waste money , value : 0.016951
feature name : return , value : 0.016890
```

```
feature name : delicious , value : 0.016376
feature name : terrible , value : 0.015325
feature name : threw , value : 0.014951
feature name : best , value : 0.014250
feature name : not recommend , value : 0.014125
feature name : good , value : 0.013403
feature name : love , value : 0.013282
```

[5.2.2] Graphviz visualization of Decision Tree on TFIDF, SET 2

```
In [160]: # Please write all the code with proper documentation
f_name = tf_idf_vect.get_feature_names()
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn import tree
from sklearn.datasets import load_wine
from IPython.display import SVG
from graphviz import Source
from IPython.display import display
dot_data = Source(export_graphviz(model, feature_names = f_name , class_names=['Positive',
'Negative'] , max_depth = 3, out_file=None, filled=True, rounded=True, special_characters=True))
display(SVG(dot_data.pipe(format='svg')))
```



[5.3] Applying Decision Trees on AVG W2V, SET 3

```
In [66]: # Please write all the code with proper documentation
depth  =[5,10,50,100,500,1000]
split  =[5,10,50,100,500,1000]
w2v_Train_score = []
w2v_val_score = []
for i in depth:
    w2v_tr = []
    w2v_val = []
    for j in split:
```

```

model= DecisionTreeClassifier(max_depth=i,min_samples_split=j)
model.fit(sent_vectors,y_tr)
train_data = model.predict_log_proba(sent_vectors)[: ,1]
val_data = model.predict_log_proba(sent_vectors_cv)[: ,1]
train_data[np.isinf(train_data)]=0
train_data[np.isnan(train_data)]=0
val_data[np.isinf(val_data)]=0
val_data[np.isnan(val_data)]=0

w2v_tr.append(roc_auc_score(np.asarray(y_tr),np.asarray(train_data)))
w2v_val.append(roc_auc_score(np.asarray(y_cv),np.asarray(val_data)))
print(str(w2v_tr)+" "+str(w2v_val))
w2v_Train_score.append(w2v_tr)
w2v_val_score.append(w2v_val)

```

```

[0.8275418763673625, 0.8275418763673625, 0.8275418763673625, 0.8275418763673625, 0.8273235388188277, 0.8261567955216202]
[0.8168153436657395, 0.8168153436657395, 0.8168153436657395, 0.8168153436657395, 0.8166493890023503, 0.8160535121348265]
[0.8378795513186862, 0.8504692081119741, 0.8927934007222851, 0.890689923691345, 0.8692552571869994, 0.8560403851508174]
[0.7757519624205393, 0.7821005737868509, 0.8230112532787663, 0.8307088844136539, 0.8357469791063745, 0.8331793148432447]
[0.5357958132796573, 0.598258696731389, 0.8342383335481327, 0.8650936134633618, 0.8803271856894314, 0.8618420111628007]
[0.5098973025946961, 0.5359097131109738, 0.7116180142499395, 0.7719182786335814, 0.8342427646443006, 0.8331255791503385]
[0.5358106321357885, 0.5987687125371948, 0.8342077867484148, 0.8653449829289868, 0.8803370520542018, 0.8618420111628007]
[0.510243766293014, 0.5379414868804623, 0.70998633905424, 0.7721487014641037, 0.834265173610556, 0.8331306108568719]
[0.5361689388843152, 0.5971752836676416, 0.8337494094062725, 0.8650409396211041, 0.8803258079679147, 0.8618424056921441]
[0.5108900427128645, 0.5366069573268759, 0.7103551017361569, 0.7719226412789546, 0.8342484562467728, 0.8331312707528107]
[0.5363667901313891, 0.5977502110314497, 0.8339304587132139, 0.8652983094814253, 0.8803378953867664, 0.8618551151731352]
[0.508976931064484, 0.5362798322456538, 0.7109761188059052, 0.7722437814706206, 0.834216341311084, 0.8331507285177843]

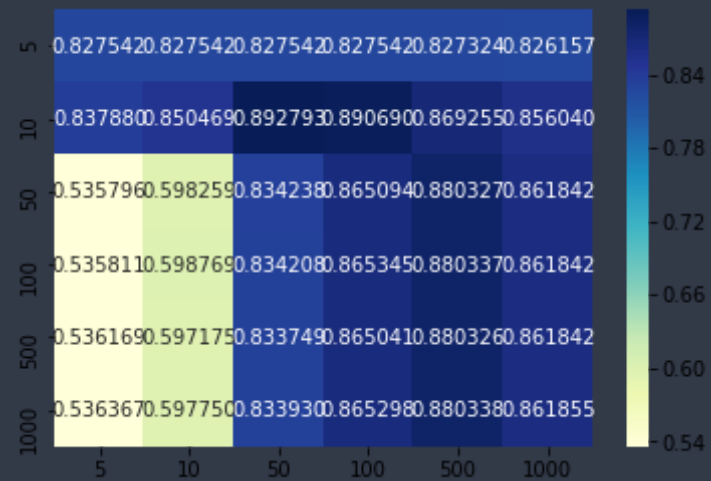
```

```

In [67]: df = pd.DataFrame(w2v_Train_score,index = depth,columns=split)
df2 = pd.DataFrame(w2v_val_score,index = depth,columns=split)
sns.heatmap(df, annot= True, fmt="f", cmap="YlGnBu")

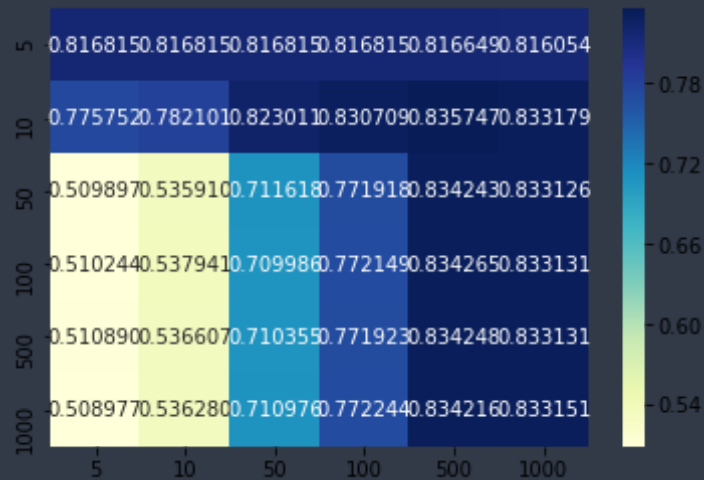
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7edde3c8>



```
In [68]: sns.heatmap(df2, annot=True, fmt="f", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7edbae80>



```
In [69]: best_depth = 10
         best_split = 1000
```


Testing on test data

```
In [70]: model= DecisionTreeClassifier(max_depth=best_depth,min_samples_split=best_split)
         model.fit(sent_vectors,y_tr)
         train_pred = model.predict_log_proba(sent_vectors)[: ,1]
         test_pred = model.predict_log_proba(sent_vectors_test)[: ,1]
         train_pred[np.isinf(train_pred)]=0
         train_pred[np.isnan(train_pred)]=0
         test_pred[np.isinf(test_pred)]=0
         test_pred[np.isnan(test_pred)]=0
         train_fpr, train_tpr, thresholds = roc_curve(y_tr, train_pred)
         test_fpr, test_tpr, thresholds = roc_curve(y_test, test_pred)

         plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
```



```
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```




Confusion matrix

```
In [71]: ytrain = model.predict(sent_vectors)
ytest = model.predict(sent_vectors_test)
ctrain = confusion_matrix(y_tr, ytrain)
ctest = confusion_matrix(y_test, ytest)
class_label = ["No", "Yes"]
```


```
df = pd.DataFrame(ctest, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7ecd0898>



```
In [73]: df = pd.DataFrame(ctrain, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7ec23390>



[5.4] Applying Decision Trees on TFIDF W2V, SET 4

```
In [75]: # Please write all the code with proper documentation
depth = [5, 10, 50, 100, 500, 1000]
split = [5, 10, 50, 100, 500, 1000]
tf_w2v_Train_score = []
tf_w2v_val_score = []
for i in depth:
    tf_w2v_tr = []
    tf_w2v_val = []
    for j in split:
        model = DecisionTreeClassifier(max_depth=i, min_samples_split=j)
        model.fit(tfidf_sent_vectors, y_tr)
        train_data = model.predict_log_proba(tfidf_sent_vectors)[:, 1]
        val_data = model.predict_log_proba(tfidf_sent_vectors_cv)[:, 1]
        train_data[np.isinf(train_data)] = 0
        train_data[np.isnan(train_data)] = 0
```

```

val_data[np.isinf(val_data)]=0
val_data[np.isnan(val_data)]=0

tf_w2v_tr.append(roc_auc_score(np.asarray(y_tr),np.asarray(train_data)))
tf_w2v_val.append(roc_auc_score(np.asarray(y_cv),np.asarray(val_data)))
print(str(tf_w2v_tr)+" "+str(tf_w2v_val))
tf_w2v_Train_score.append(tf_w2v_tr)
tf_w2v_val_score.append(tf_w2v_val)

```

```

[0.7997598130406851, 0.7997598130406851, 0.7997598130406851, 0.7997598130406851, 0.7994837583541587, 0.7986662371932958]
[0.779696499559996, 0.779696499559996, 0.779696499559996, 0.779696499559996, 0.7799550962810171, 0.7794206263966881]
[0.8206468561167349, 0.8301246950695229, 0.8643222845912578, 0.8641294390656945, 0.8455161270444604, 0.8308300321247083]
[0.7512611802869433, 0.7562244776080445, 0.7815636659536368, 0.7938044331662526, 0.8016060895637165, 0.7993246001727168]
[0.5372022331529813, 0.6124633854851026, 0.8136409689456977, 0.8592514479957514, 0.8589524511147876, 0.8353053579882026]
[0.5087878892084577, 0.5483809123193332, 0.676731094018014, 0.739740506187111, 0.7997779028565722, 0.799709933574875]
[0.5367881746070812, 0.6113530097850374, 0.8123494834024427, 0.8592928224341159, 0.8590637783194311, 0.8353053579882026]
[0.5080393839094202, 0.546171892333725, 0.6777160620284586, 0.7401670830854447, 0.79968091648401, 0.799709933574875]
[0.5374811152998019, 0.6141775330276377, 0.8136428413944863, 0.8596136207143519, 0.8589581811838226, 0.8353053579882026]
[0.509447455199298, 0.5469530716469551, 0.6770832951449111, 0.7404159188457041, 0.7995710254797554, 0.799709933574875]
[0.5366690559697695, 0.6121785206869687, 0.8134877997123567, 0.8594202387123797, 0.8589581811838226, 0.8353053579882026]
[0.5079989011266184, 0.5462133282995494, 0.67853244495698, 0.7399443773713177, 0.7995369400214717, 0.799709933574875]

```

```


In [76]: df = pd.DataFrame(tf_w2v_Train_score,index = depth,columns=split)
df2 = pd.DataFrame(tf_w2v_val_score,index = depth,columns=split)
sns.heatmap(df, annot= True, fmt="f", cmap="YlGnBu")

```

```


<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7ebf0b70>

```



```
In [81]: sns.heatmap(df2, annot=True, fmt="f", cmap="YlGnBu")
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7eb2a400>



```
In [77]: best_depth = 10  
best_split = 500
```

Testing on test data


```
In [78]: model= DecisionTreeClassifier(max_depth=best_depth,min_samples_split=best_split)  
model.fit(tfidf_sent_vectors,y_tr)  
train_pred = model.predict_log_proba(tfidf_sent_vectors)[: ,1]  
test_pred = model.predict_log_proba(tfidf_sent_vectors_test)[: ,1]  
train_pred[np.isinf(train_pred)]=0  
train_pred[np.isnan(train_pred)]=0  
test_pred[np.isinf(test_pred)]=0  
test_pred[np.isnan(test_pred)]=0  
train_fpr, train_tpr, thresholds = roc_curve(y_tr, train_pred)  
test_fpr, test_tpr, thresholds = roc_curve(y_test, test_pred)  
  
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))  
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))  
plt.plot([0.0,1.0],[0.0,1.0])  
plt.legend()  
plt.xlabel("alpha: hyperparameter")  
plt.ylabel("AUC")  
plt.title("ERROR PLOTS")  
plt.show()
```



confusion matrix


```
In [175]: ytrain = model.predict(tfidf_sent_vectors)
ytest = model.predict(tfidf_sent_vectors_test)
ctrain = confusion_matrix(y_tr,ytrain)
ctest = confusion_matrix(y_test,ytest)
class_label=["No","Yes"]
df = pd.DataFrame(ctest, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fcbb464ff60>



```
In [79]: df = pd.DataFrame(ctrain, index=class_label, columns=class_label)
sns.heatmap(df, annot=True, fmt="d", cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8e7eb6dbe0>



[6] Conclusions

```
In [74]: # Please compare all your models using Prettytable library
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "value of depth", "value of split", "Train AUC", "Test AUC"]
x.add_row(["BOW", 50, 500, 0.84, 0.82])
x.add_row(["TFIDF", 50, 500, 0.84, 0.81])
x.add_row(["Avg_w2v", 5, 500, 0.85, 0.83])
x.add_row(["TFIDF_w2v", 10, 500, 0.84, 0.80])

print(x)
```

Model	value of depth	value of split	Train AUC	Test AUC
BOW	50	500	0.84	0.82
TFIDF	50	500	0.84	0.81
Avg_w2v	5	500	0.85	0.83
TFIDF_w2v	10	500	0.84	0.8

Here we can see that our Avg_W2V model got the highest AUC value and is the best model with min_spli = 500 and depth = 5